

CPSC 340: Machine Learning and Data Mining

Non-Parametric Models

Term2, 2021

Admin

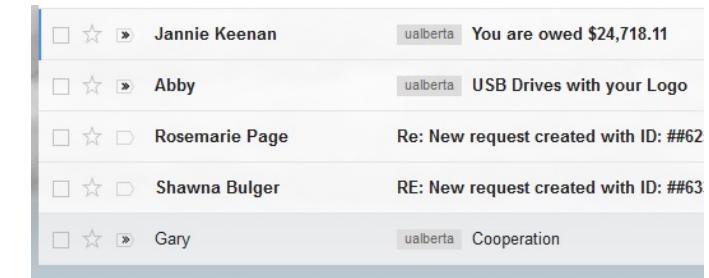
- **Add/drop deadline** is today.
- **Audit:** Anyone who wants to audit the course should email their registration form to cpsc340-admin@cs.ubc.ca.

Last Time: E-mail Spam Filtering

- Want a build a system that filters spam e-mails:

- We formulated as **supervised learning**:

- $(y_i = 1)$ if e-mail ‘i’ is spam, $(y_i = 0)$ if e-mail is not spam.
 - $(x_{ij} = 1)$ if word/phrase ‘j’ is in e-mail ‘i’, $(x_{ij} = 0)$ if it is not.



\$	Hi	CPSC	340	Vicodin	Offer	...		Spam?
1	1	0	0	1	0	...		1
0	0	0	0	1	1	...		1
0	1	1	1	0	0	...		0
...		...

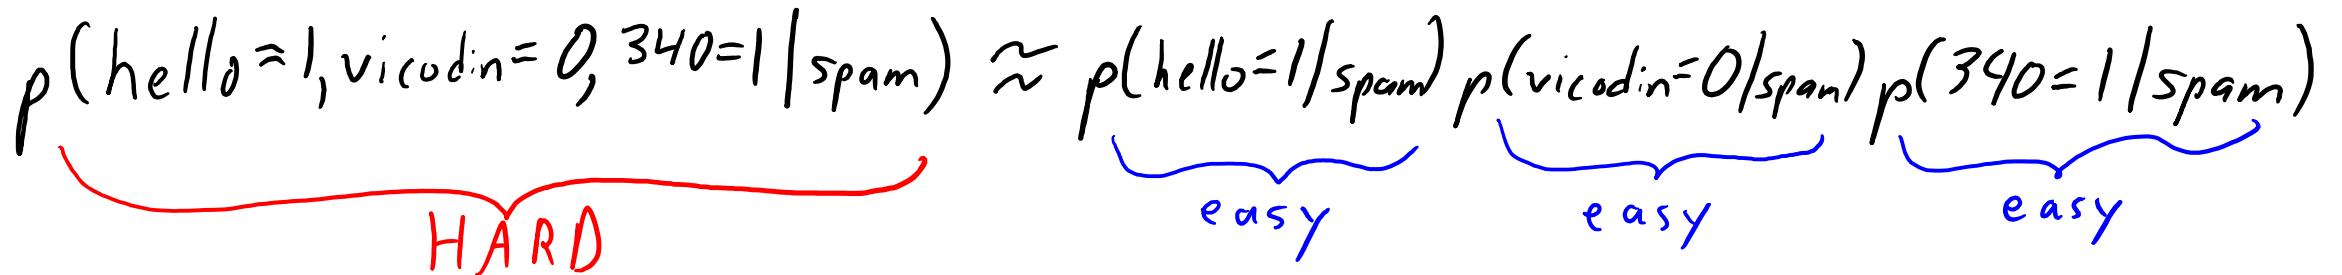
Last Time: Naïve Bayes

- We considered spam filtering methods based on naïve Bayes:

$$p(y_i = \text{"spam"} | x_i) = \frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- Makes conditional independence assumption to make learning practical:

$$p(\text{hello}=1, \text{vicodin}=0, \text{340}=1 | \text{spam}) \approx p(\text{hello}=1 | \text{spam}) p(\text{vicodin}=0 | \text{spam}) p(\text{340}=1 | \text{spam})$$



HARD

- Predict “spam” if $p(y_i = \text{"spam"} | x_i) > p(y_i = \text{"not spam"} | x_i)$.
 - We don’t need $p(x_i)$ to test this.

Naïve Bayes

- Naïve Bayes formally:

$$\begin{aligned} p(y_i | x_i) &= \frac{p(x_i | y_i) p(y_i)}{p(x_i)} \quad (\text{first use Bayes rule}) \\ &\propto p(x_i | y_i) p(y_i) \quad ("denominator doesn't matter") \\ &\approx \prod_{j=1}^d \left[p(x_{ij} | y_i) \right] p(y_i) \quad (\text{conditional independence assumption}) \end{aligned}$$

Only needs easy probabilities.

same for all y_i values

- Post-lecture slides: how to train/test by hand on a simple example.

Laplace Smoothing

- Our estimate of $p(\text{'lactase'} = 1 | \text{'spam'})$ is:

$$\frac{\text{\# spam messages with lactase}}{\text{\# spam messages}}$$

- But there is a problem if you have **no spam messages with lactase**:
 - $p(\text{'lactase'} | \text{'spam'}) = 0$, so spam messages with lactase automatically get through.
- Common fix is **Laplace smoothing**:
 - Add 1 to numerator, and 2 to denominator (for binary features).
 - Acts like a “fake” spam example that has lactase, and a “fake” spam example that doesn’t.

Laplace Smoothing

- Laplace smoothing:

$$\frac{(\text{\# Spam messages with lactase}) + 1}{(\text{\# spam messages}) + 2}$$

- Typically you do this for all features.
 - Helps against overfitting by biasing towards the uniform distribution.
- A common variation is to use a real number β rather than 1.
 - Add ' βk ' to denominator if feature has 'k' possible values (so it sums to 1).

$$p(x_{ij}=c | y_i = \text{class}) \approx \frac{(\text{number of examples in class with } x_{ij}=c) + \beta}{(\text{number of examples in class}) + \beta K}$$

This is a “maximum a posteriori” (MAP) estimate of the probability. We’ll discuss MAP and how to derive this formula later. bonus!

Decision Theory

- Are we equally concerned about “spam” vs. “not spam”?
- True positives, false positives, false negatives, true negatives:

Predict / True	True ‘spam’	True ‘not spam’
Predict ‘spam’	True Positive	False Positive
Predict ‘not spam’	False Negative	True Negative

- The costs mistakes might be different:
 - Letting a spam message through (false negative) is not a big deal.
 - Filtering a not spam (false positive) message will make users mad.

Decision Theory

- We can give a **cost** to each scenario, such as:

Predict / True	True 'spam'	True 'not spam'
Predict 'spam'	0	100
Predict 'not spam'	10	0

- Instead of most probable label, take \hat{y}_i minimizing **expected cost**:

$$E \left[\underbrace{\text{cost}(\hat{y}_i, \tilde{y}_i)}_{\substack{\text{expectation of model} \\ \text{with respect to } \hat{y}_i}} \right] \xrightarrow{\text{Cost of predicting } \hat{y}_i \text{ if it's really } \tilde{y}_i}$$

- Even if “spam” has a higher probability, predicting “spam” might have a higher expected cost.

Decision Theory Example

Predict / True	True 'spam'	True 'not spam'
Predict 'spam'	0	100
Predict 'not spam'	10	0

- Consider a test example we have $p(\tilde{y}_i = \text{"spam"} \mid \tilde{x}_i) = 0.6$, then:

$$\begin{aligned} \mathbb{E}[\text{cost}(\hat{y}_i = \text{"spam"}, \tilde{y}_i)] &= p(\tilde{y}_i = \text{"spam"} \mid \tilde{x}_i) \text{cost}(\hat{y}_i = \text{"spam"}, \tilde{y}_i = \text{"spam"}) \\ &\quad + p(\tilde{y}_i = \text{"not spam"} \mid \tilde{x}_i) \text{cost}(\hat{y}_i = \text{"spam"}, \tilde{y}_i = \text{"not spam"}) \\ &= (0.6)(0) + (0.4)(100) = 40 \end{aligned}$$

$$\mathbb{E}[\text{cost}(\hat{y}_i = \text{"not spam"}, \tilde{y}_i)] = (0.6)(10) + (0.4)(0) = 6$$

- Even though "spam" is more likely, we should predict "not spam".

bonus!

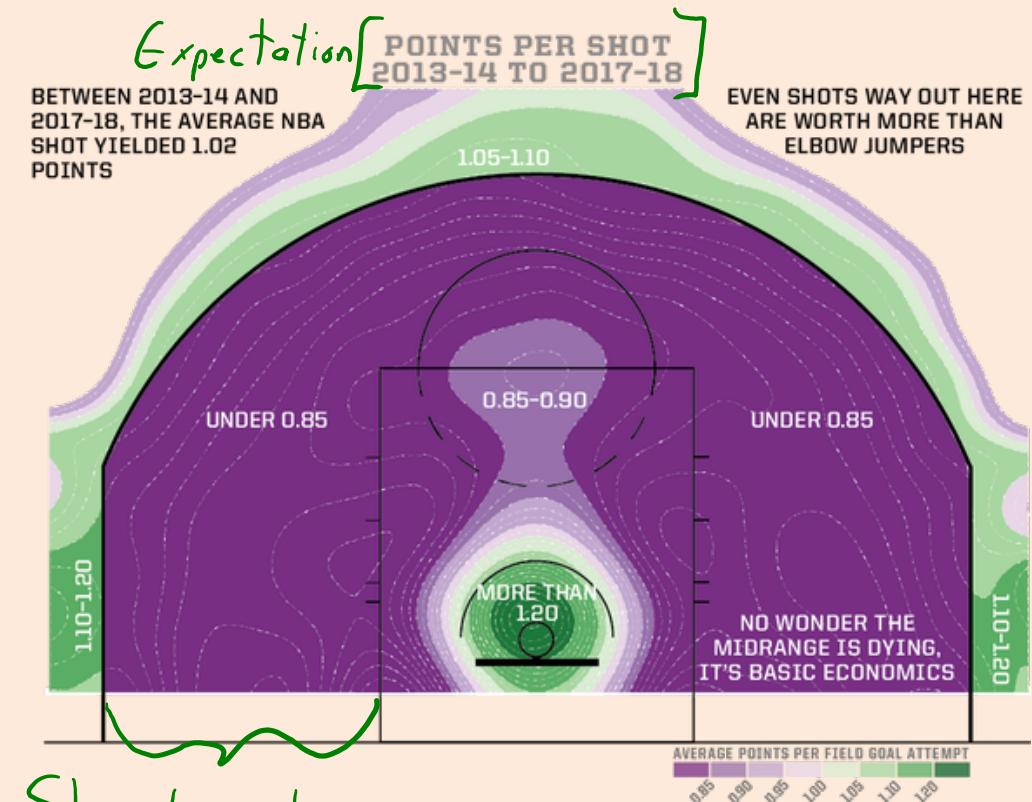
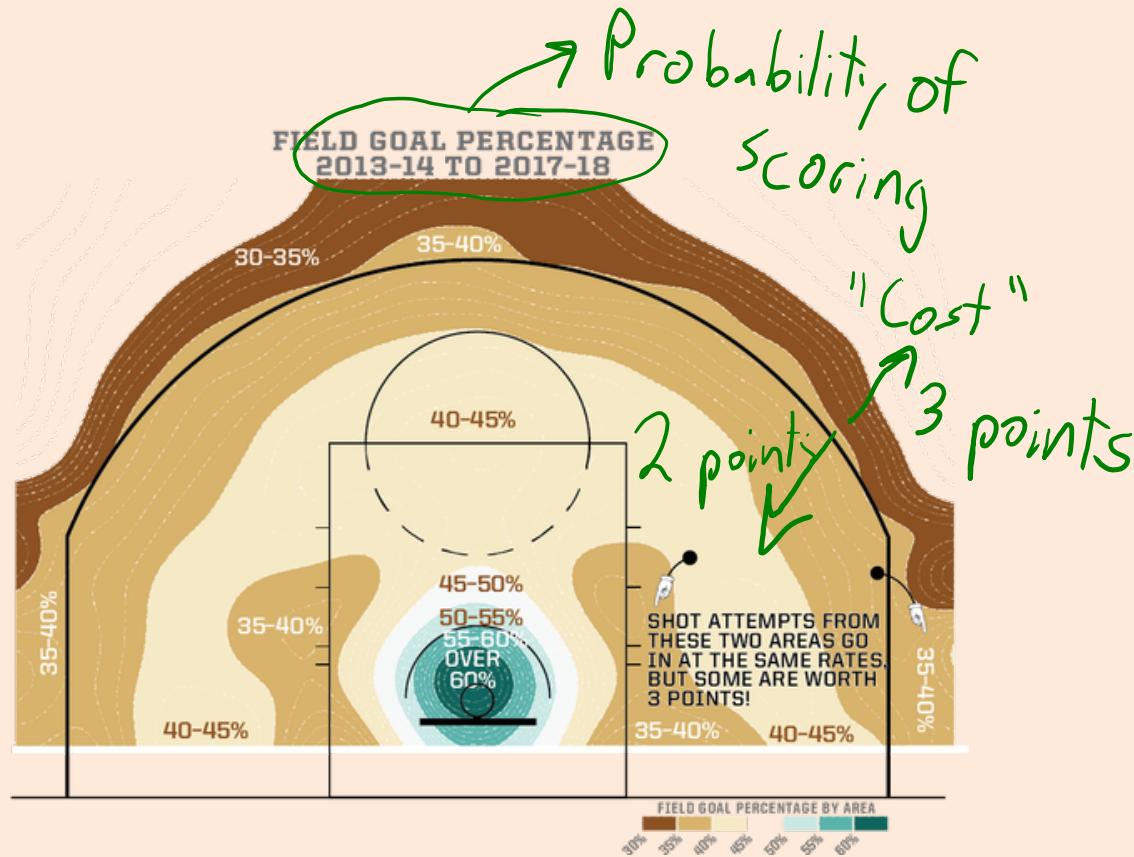
Decision Theory Discussion

- In other applications, the costs could be different.
 - In cancer screening, some false positives are ok, but don't want to have false negatives.
- Decision theory and “darts”:
 - <http://www.datagenetics.com/blog/january12012/index.html>
- Decision theory and video poker:
 - <http://datagenetics.com/blog/july32019/index.html>
- Decision theory can help with “unbalanced” class labels:
 - If 99% of e-mails are spam, you get 99% accuracy by always predicting “spam”.
 - Decision theory approach avoids this.
 - See also [precision/recall curves](#) and [ROC curves](#) in the bonus material.

bonus!

Decision Theory and Basketball

- “How Mapping Shots In The NBA Changed It Forever”

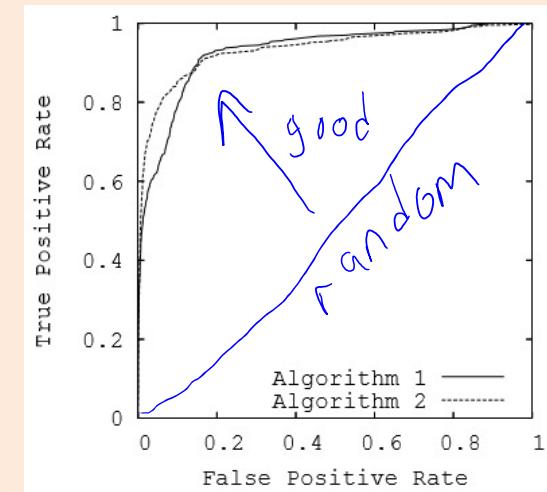


Shooting here is
a bad decision

bonus!

Unbalanced Class Labels

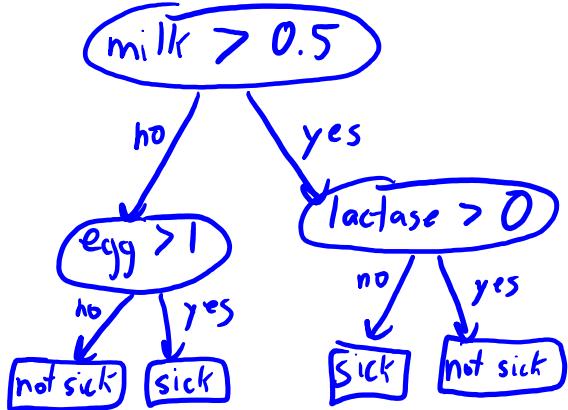
- A related is that of “**unbalanced**” class labels.
 - If 99% of the e-mails are spam, you can get 99% accuracy by always predicting spam.
- There are a variety of other performance measures available:
 - Weighted classification error.
 - Jaccard similarity.
 - Precision and recall.
 - False positive and false negative rate.
 - ROC curves.
- See the post-lecture bonus slides for additional details.



(pause)

Decision Trees vs. Naïve Bayes

- Decision trees:



1. Sequence of rules based on 1 feature.
2. Training: 1 pass over data per depth.
3. Greedy splitting as approximation.
4. Testing: just look at features in rules.
5. New data: might need to change tree.
6. Accuracy: good if simple rules based on individual features work (“symptoms”).

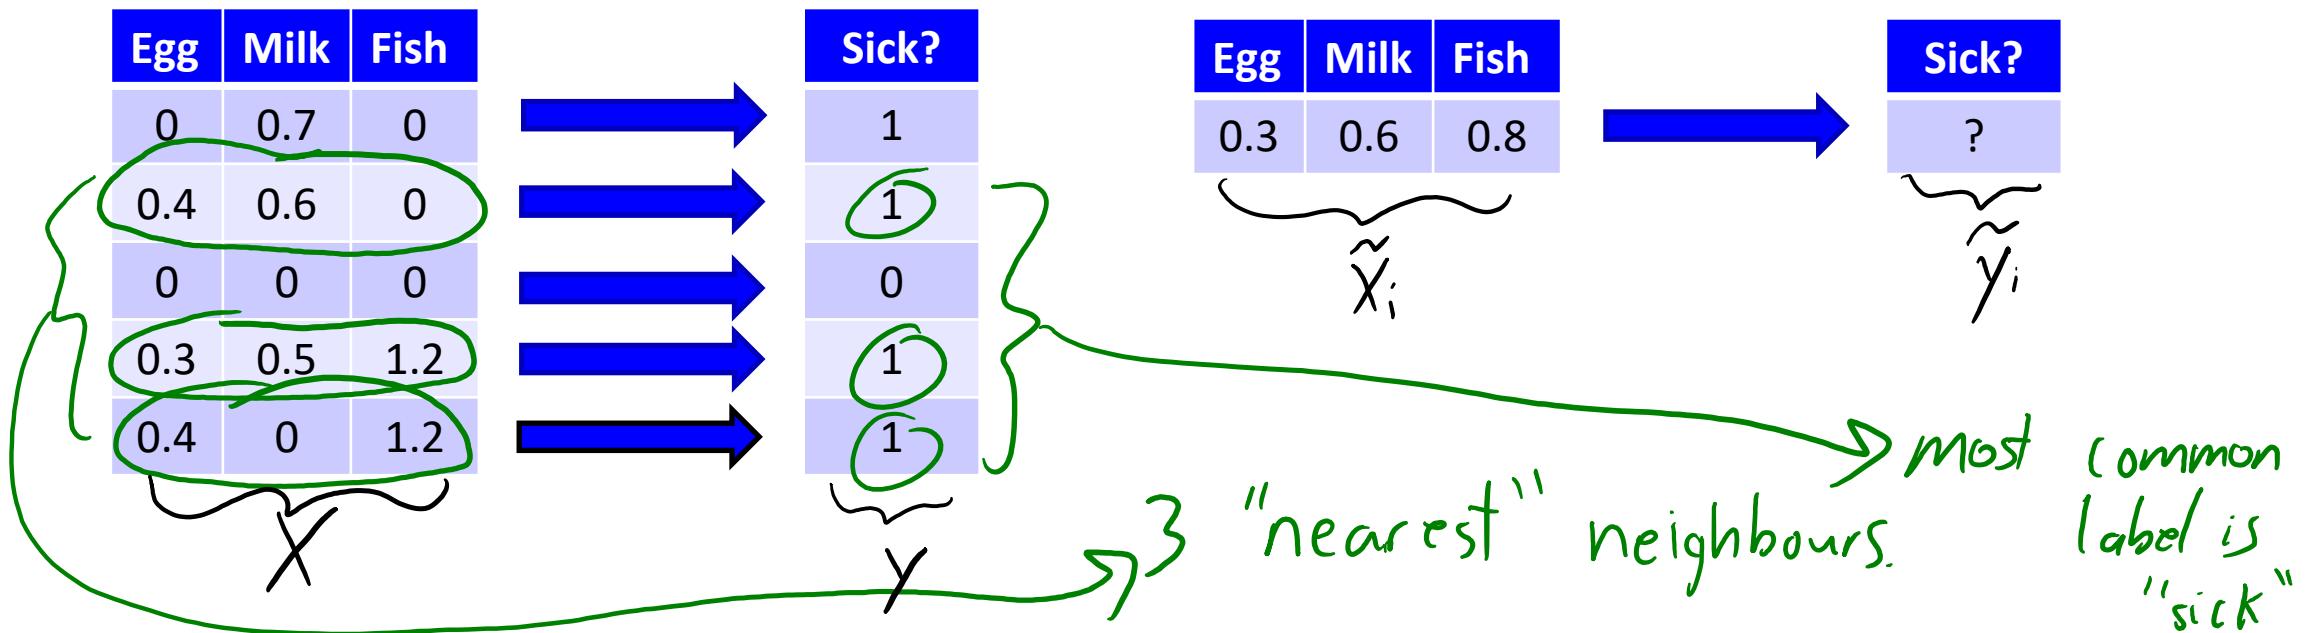
- Naïve Bayes:

$$\begin{aligned} & p(\text{sick} \mid \text{milk, egg, lactase}) \\ & \approx p(\text{milk} \mid \text{sick}) p(\text{egg} \mid \text{sick}) p(\text{lactase} \mid \text{sick}) p(\text{sick}) \end{aligned}$$

1. Simultaneously combine all features.
2. Training: 1 pass over data to count.
3. Conditional independence assumption.
4. Testing: look at all features.
5. New data: just update counts.
6. Accuracy: good if features almost independent given label (bag of words).

k-Nearest Neighbours (kNN)

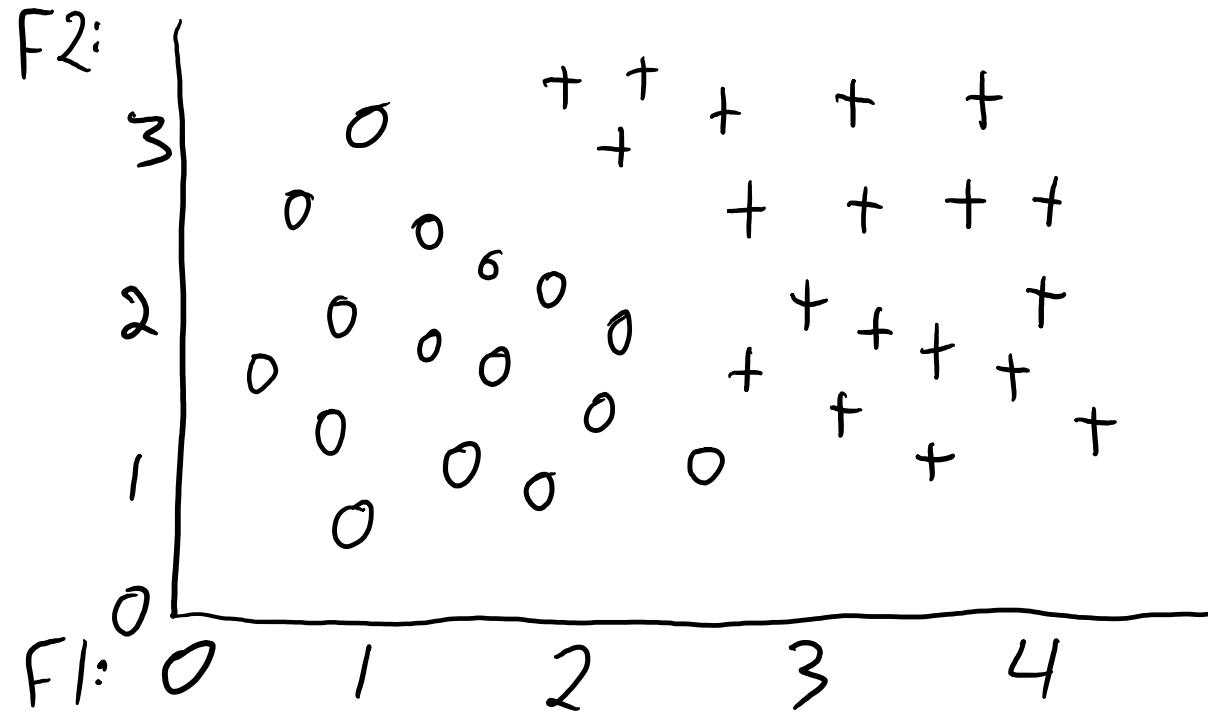
- An old/simple classifier: k-nearest neighbours (kNN).
- To classify an example \tilde{x}_i :
 1. Find the 'k' training examples x_i that are “nearest” to \tilde{x}_i .
 2. Classify using the most common label of “nearest” training examples.



k-Nearest Neighbours (kNN)

- An old/simple classifier: **k-nearest neighbours (kNN)**.
- To classify an example \tilde{x}_i :
 1. Find **the 'k'** training examples x_i that are “nearest” to \tilde{x}_i .
 2. Classify using the **most common label** of “nearest” training examples.

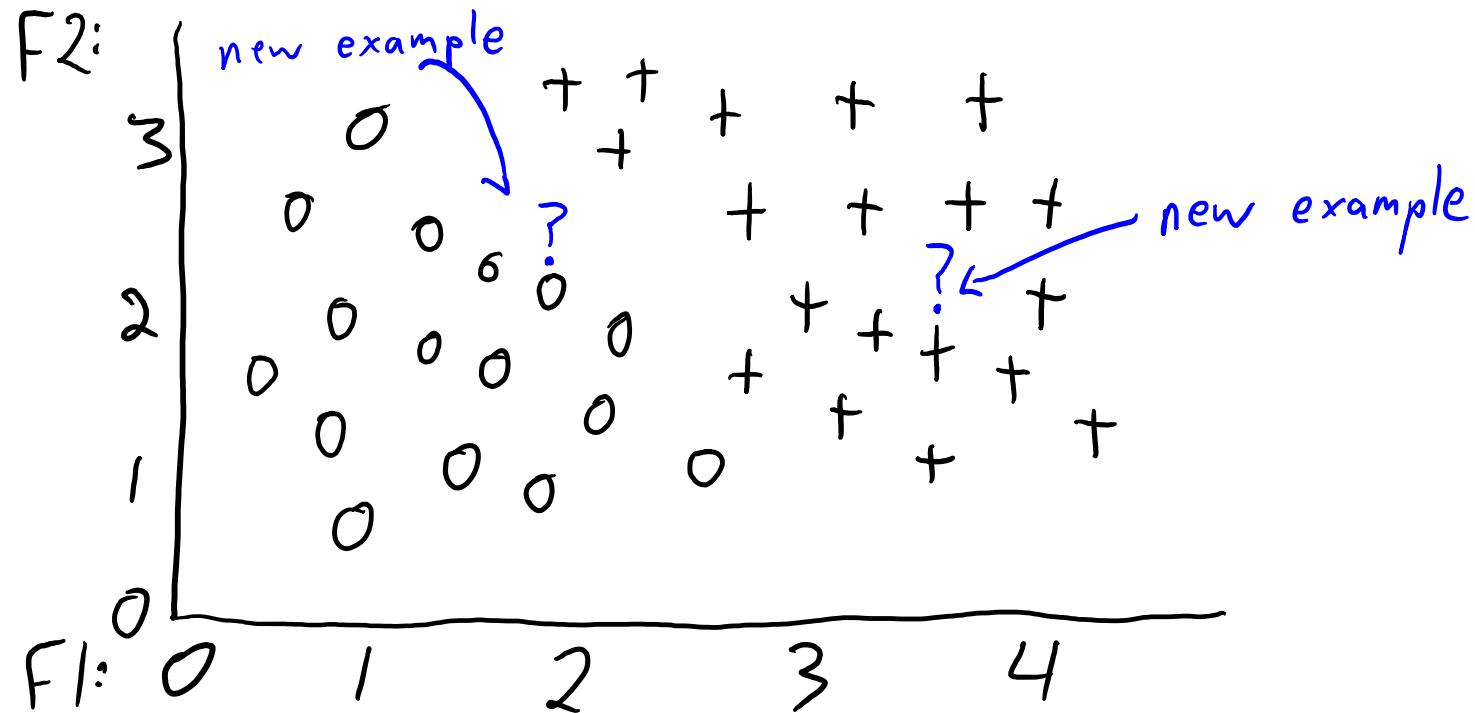
F1	F2	
Label		
1	3	0
2	3	+
3	2	+
2.5	1	0
3.5	1	+
...



k-Nearest Neighbours (kNN)

- An old/simple classifier: k-nearest neighbours (kNN).
- To classify an example \tilde{x}_i :
 1. Find the 'k' training examples x_i that are “nearest” to \tilde{x}_i .
 2. Classify using the most common label of “nearest” training examples.

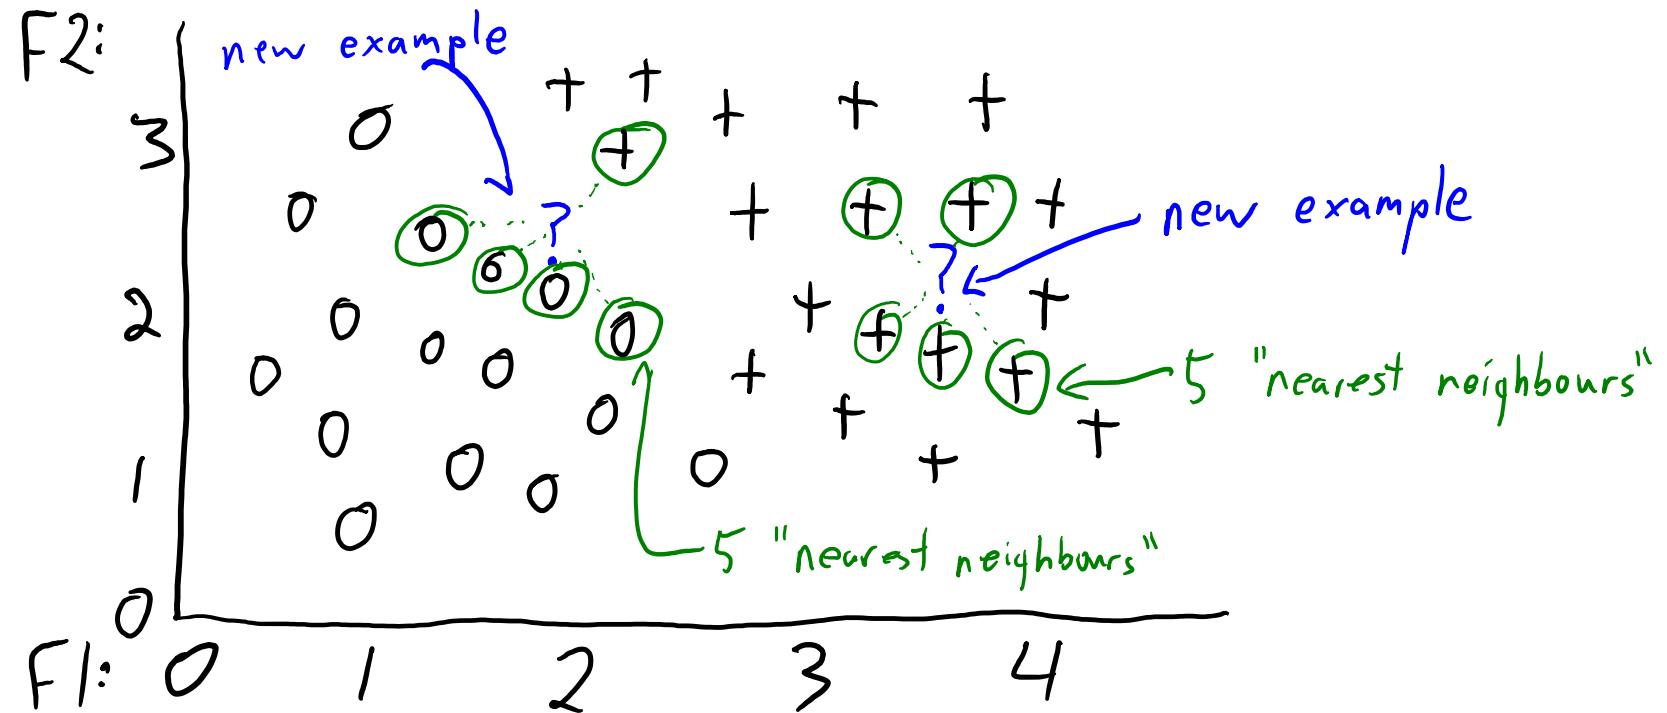
F1	F2	
Label		
1	3	0
2	3	+
3	2	+
2.5	1	0
3.5	1	+
...



k-Nearest Neighbours (kNN)

- An old/simple classifier: k-nearest neighbours (kNN).
- To classify an example \tilde{x}_i :
 1. Find the 'k' training examples x_i that are “nearest” to \tilde{x}_i .
 2. Classify using the most common label of “nearest” training examples.

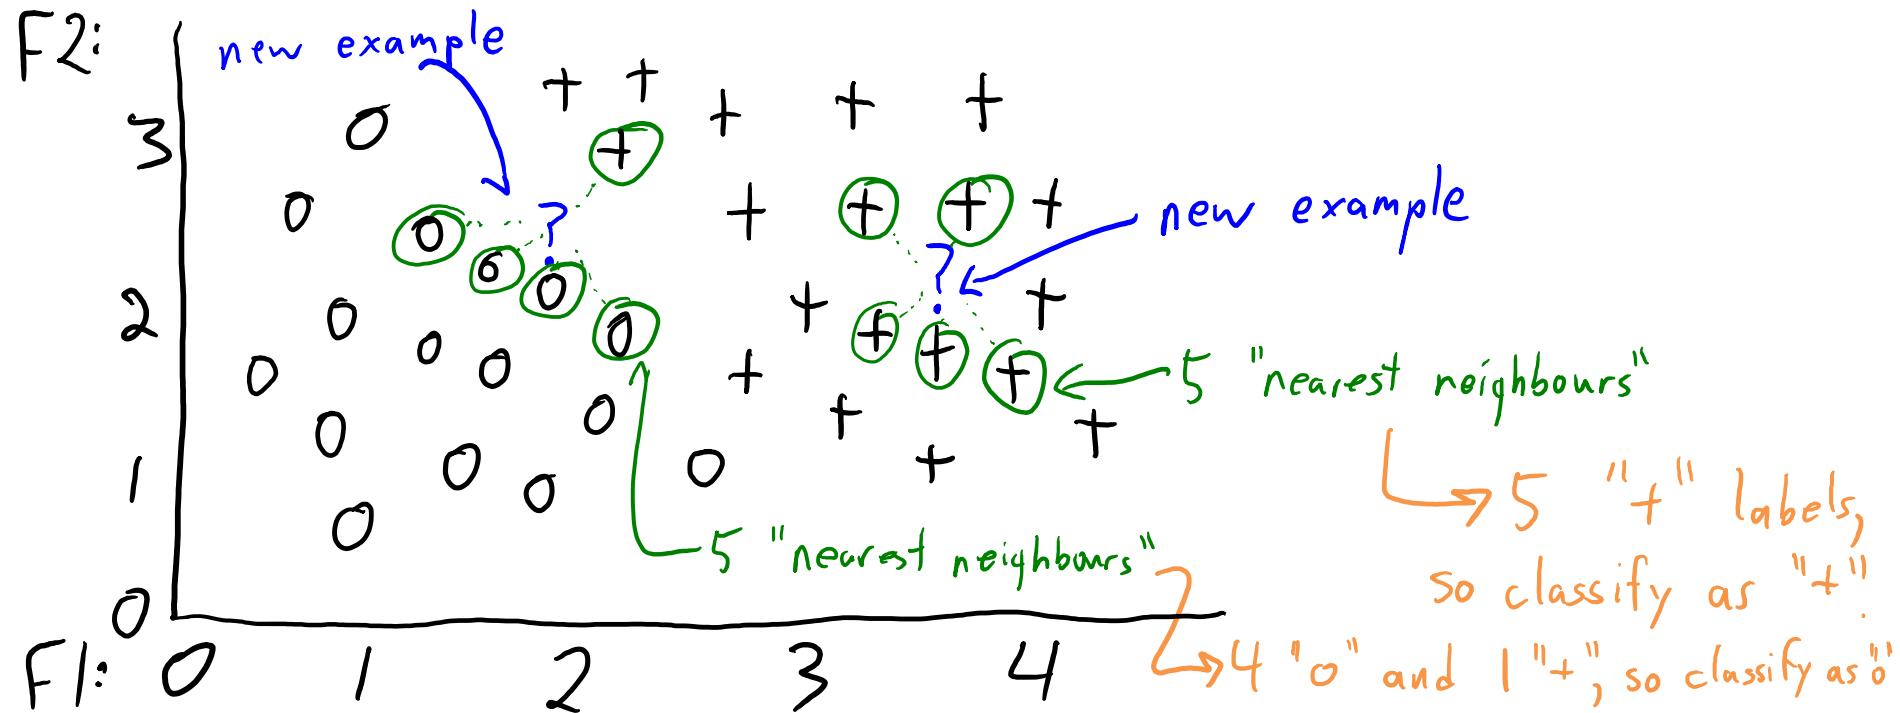
F1	F2	
1	3	→
2	3	→
3	2	→
2.5	1	→
3.5	1	→
...



k-Nearest Neighbours (kNN)

- An old/simple classifier: k-nearest neighbours (kNN).
- To classify an example \tilde{x}_i :
 1. Find the 'k' training examples x_i that are “nearest” to \tilde{x}_i .
 2. Classify using the most common label of “nearest” training examples.

F1	F2	Label
1	3	0
2	3	+
3	2	+
2.5	1	0
3.5	1	+
...



k-Nearest Neighbours (kNN)

- Assumption:
 - Examples with similar features are likely to have similar labels.
- Seems strong, but all good classifiers basically rely on this assumption.
 - If not true there may be nothing to learn and you are in “no free lunch” territory.
 - Methods just differ in how you define “similarity”.
- Most common distance function is Euclidean distance:

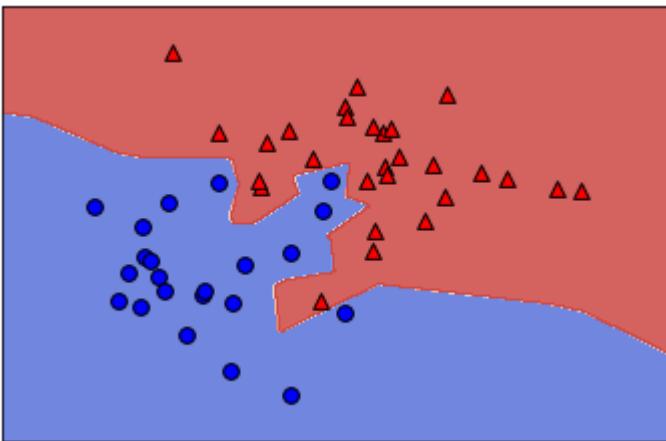
$$\|x_i - \tilde{x}_{\tilde{i}}\| = \sqrt{\sum_{j=1}^d (x_{ij} - \tilde{x}_{\tilde{i}j})^2}$$

- x_i is features of training example ‘i’, and $\tilde{x}_{\tilde{i}}$ is features of test example ‘ \tilde{i} ’.
- Costs $O(d)$ to calculate for a pair of examples.

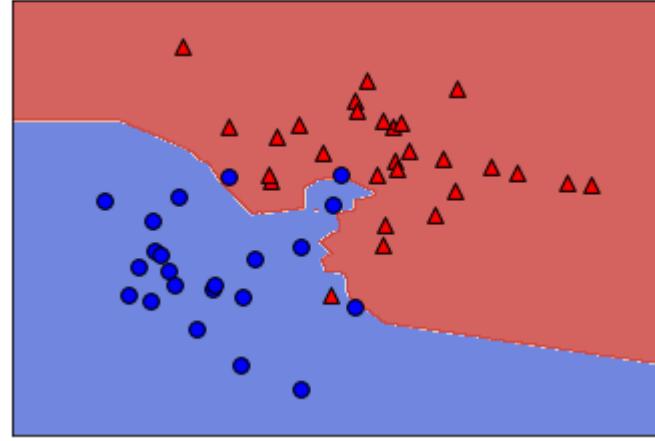
Effect of 'k' in kNN.

- With large 'k' (hyper-parameter), kNN model will be very simple.
 - With $k=n$, you just predict the mode of the labels.
 - Model **gets more simple** as 'k' increases.

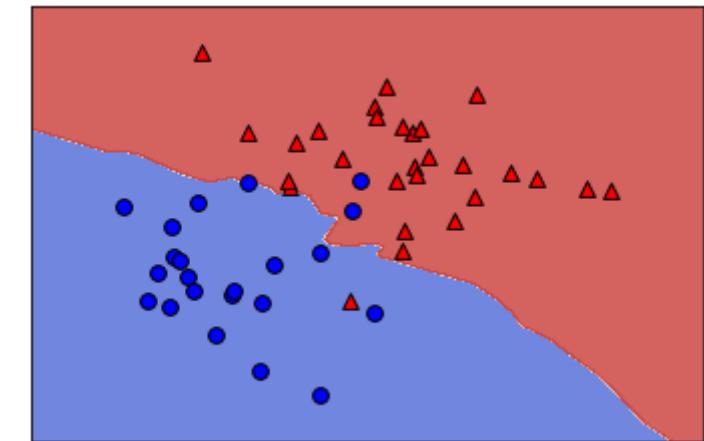
$k=1$



$k=3$



$k=10$



- Effect of 'k' on fundamental trade-off:
 - As 'k' grows, training error increases and approximation error decreases.

kNN Implementation

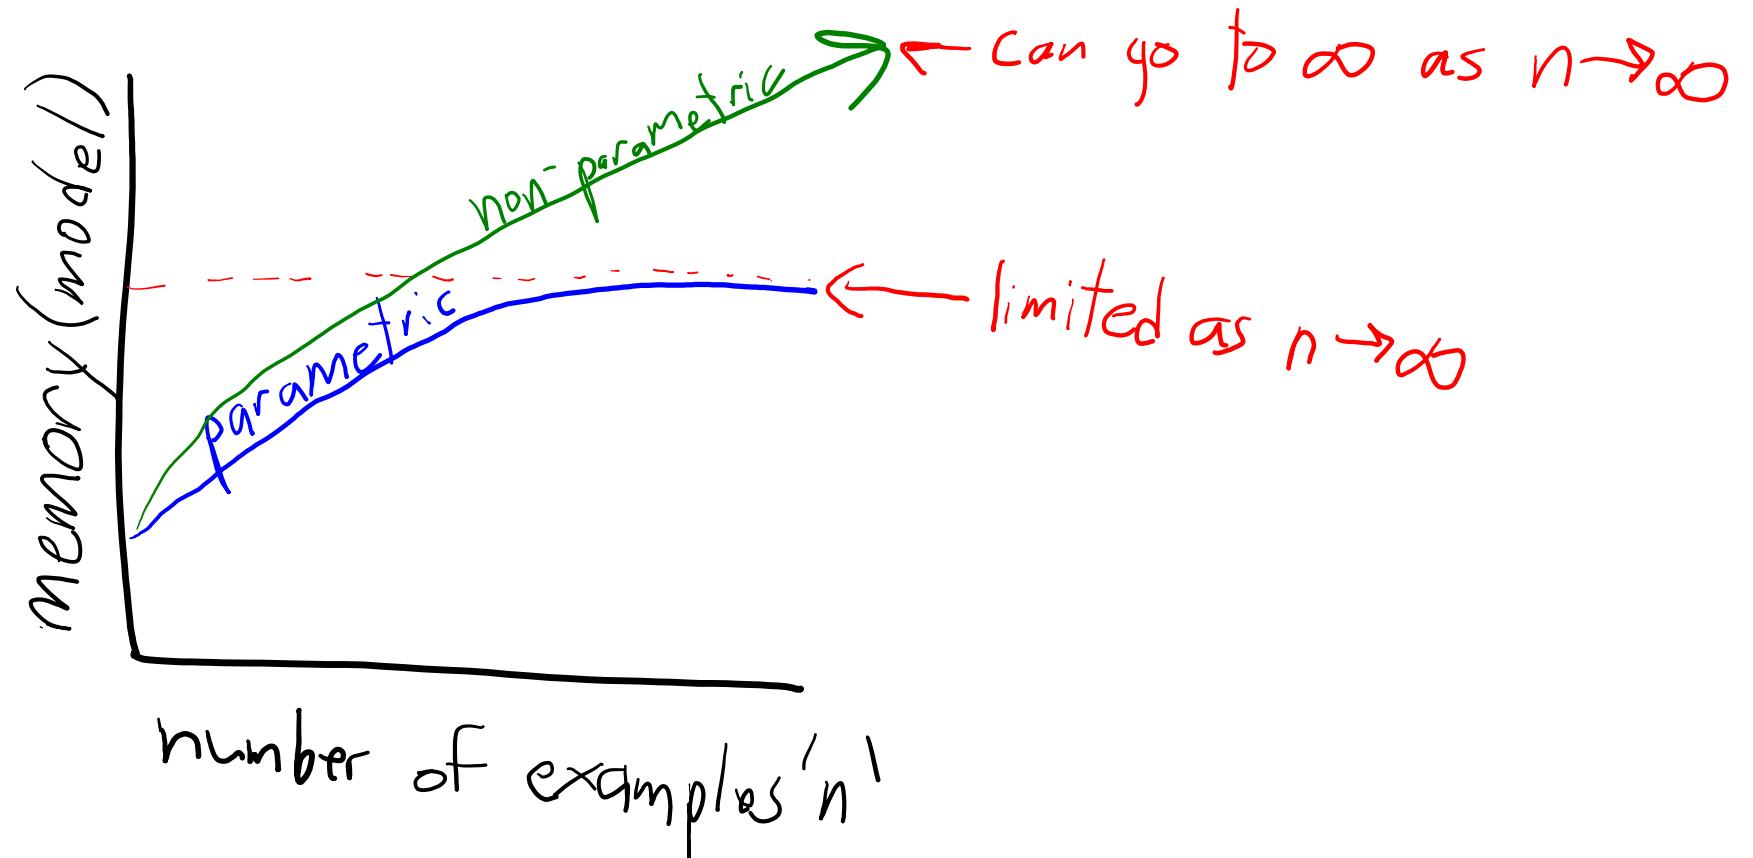
- There is **no training** phase in kNN (“lazy” learning).
 - You just store the training data.
 - Costs $O(1)$ if you use a pointer.
- But **predictions are expensive**: $O(nd)$ to classify 1 test example.
 - Need to do $O(d)$ distance calculation for all ‘n’ training examples.
 - So **prediction time grows with number of training examples**.
 - Tons of work on reducing this cost (we’ll discuss this later).
- But **storage is expensive**: needs $O(nd)$ memory to store ‘X’ and ‘y’.
 - So **memory grows with number of training examples**.
 - When storage depends on ‘n’, we call it a **non-parametric** model.

Parametric vs. Non-Parametric

- **Parametric** models:
 - Have **fixed number** of parameters: trained “model” size is $O(1)$ in terms ‘n’.
 - E.g., naïve Bayes just stores counts.
 - E.g., fixed-depth decision tree just stores rules for that depth.
 - You can estimate the fixed parameters more accurately with more data.
 - But **eventually more data doesn’t help**: model is too simple.
- **Non-parametric** models:
 - **Number of parameters grows with ‘n’**: size of “model” depends on ‘n’.
 - Model gets **more complicated as you get more data**.
 - E.g., kNN stores all the training data, so size of “model” is $O(nd)$.
 - E.g., decision tree whose depth *grows with the number of examples*.

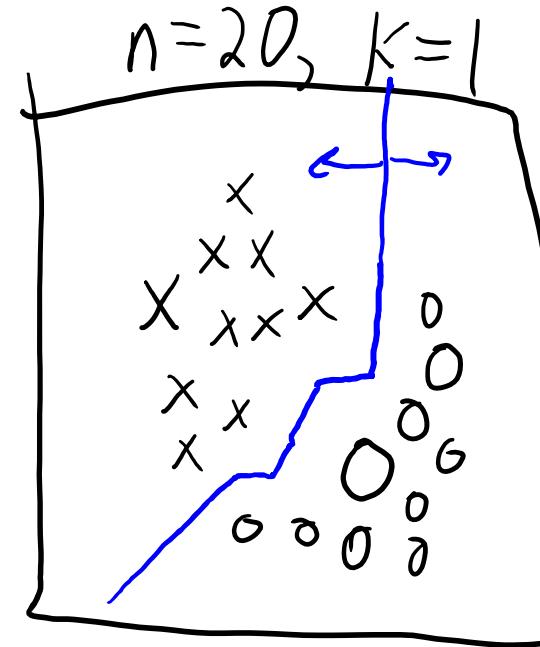
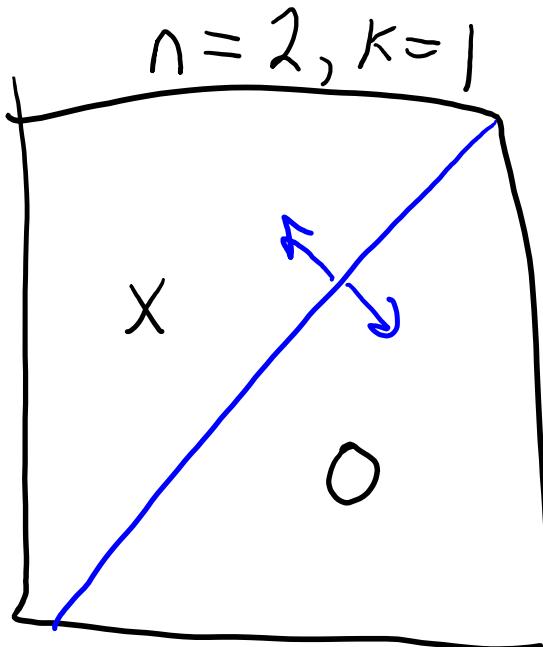
Parametric vs. Non-Parametric Models

- Parametric models have bounded memory.
- Non-parametric models can have unbounded memory.



Effect of 'n' in kNN.

- With a small 'n', kNN model will be very simple.



- Model gets more complicated as 'n' increases.
 - Requires more memory, but detects subtle differences between examples.

bonus!

Consistency of kNN ($n \rightarrow \infty$)

- KNN has appealing **consistency** properties:
 - As 'n' goes to ∞ , KNN test error is **at most twice the best possible error**.
 - For fixed 'k' and binary labels (under mild assumptions).
- Stone's Theorem: kNN is “**universally consistent**”.
 - If k/n goes to zero and 'k' goes to ∞ , **converges to the best possible error**.
 - For example, $k = \log(n)$.
 - First algorithm shown to have this property.
- Does Stone's Theorem violate the no free lunch theorem?
 - No: it requires a continuity assumption on the labels.
 - Consistency says nothing about finite 'n' (see "[Dont Trust Asymptotics](#)").
 - The “speed” at which universal consistency happens is **exponential in the dimension 'd'**.

Curse of Dimensionality

- “Curse of dimensionality”: problems with high-dimensional spaces.
 - Volume of space grows **exponentially** with dimension.
 - Circle has area $O(r^2)$, sphere has area $O(r^3)$, 4d hyper-sphere has area $O(r^4)$,...
 - Need **exponentially more points** to ‘fill’ a high-dimensional volume.
 - “Nearest” neighbours might be **really far** even with large ‘n’.
- KNN is also problematic if features have very **different scales**.
 - Comparing a feature measured in grams vs one measure in kilograms.
 - Measurement in grams can have much more influence (values 1000 times larger).
- Nevertheless, KNN is **really easy to use and often hard to beat!**

bonus!

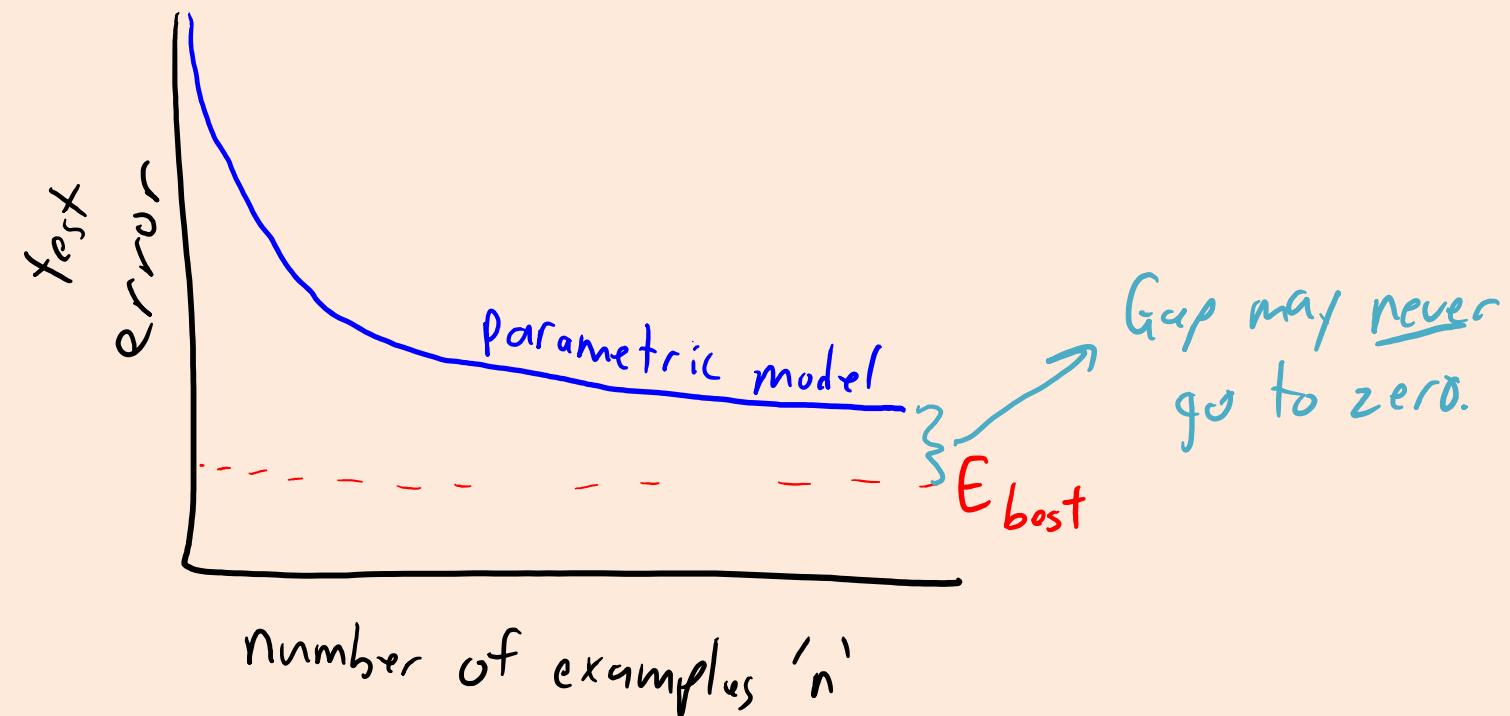
Consistency of Non-Parametric Models

- **Universal consistency** can be shown for many models in 340:
 - “Linear” models with “polynomial” or “RBFs” as features (later).
 - “Neural network” and “deep learning” models (also covered later).
- But it's always the **non-parametric versions** that are consistent:
 - Where **size of model** is a function of ‘n’.
 - Examples:
 - KNN needs to store all ‘n’ training examples.
 - Degree of the polynomial must grow with ‘n’ (not true for fixed polynomial).
 - Number of “hidden units” must grow with ‘n’ (not true for fixed neural network).

bonus!

Parametric vs. Non-Parametric Models

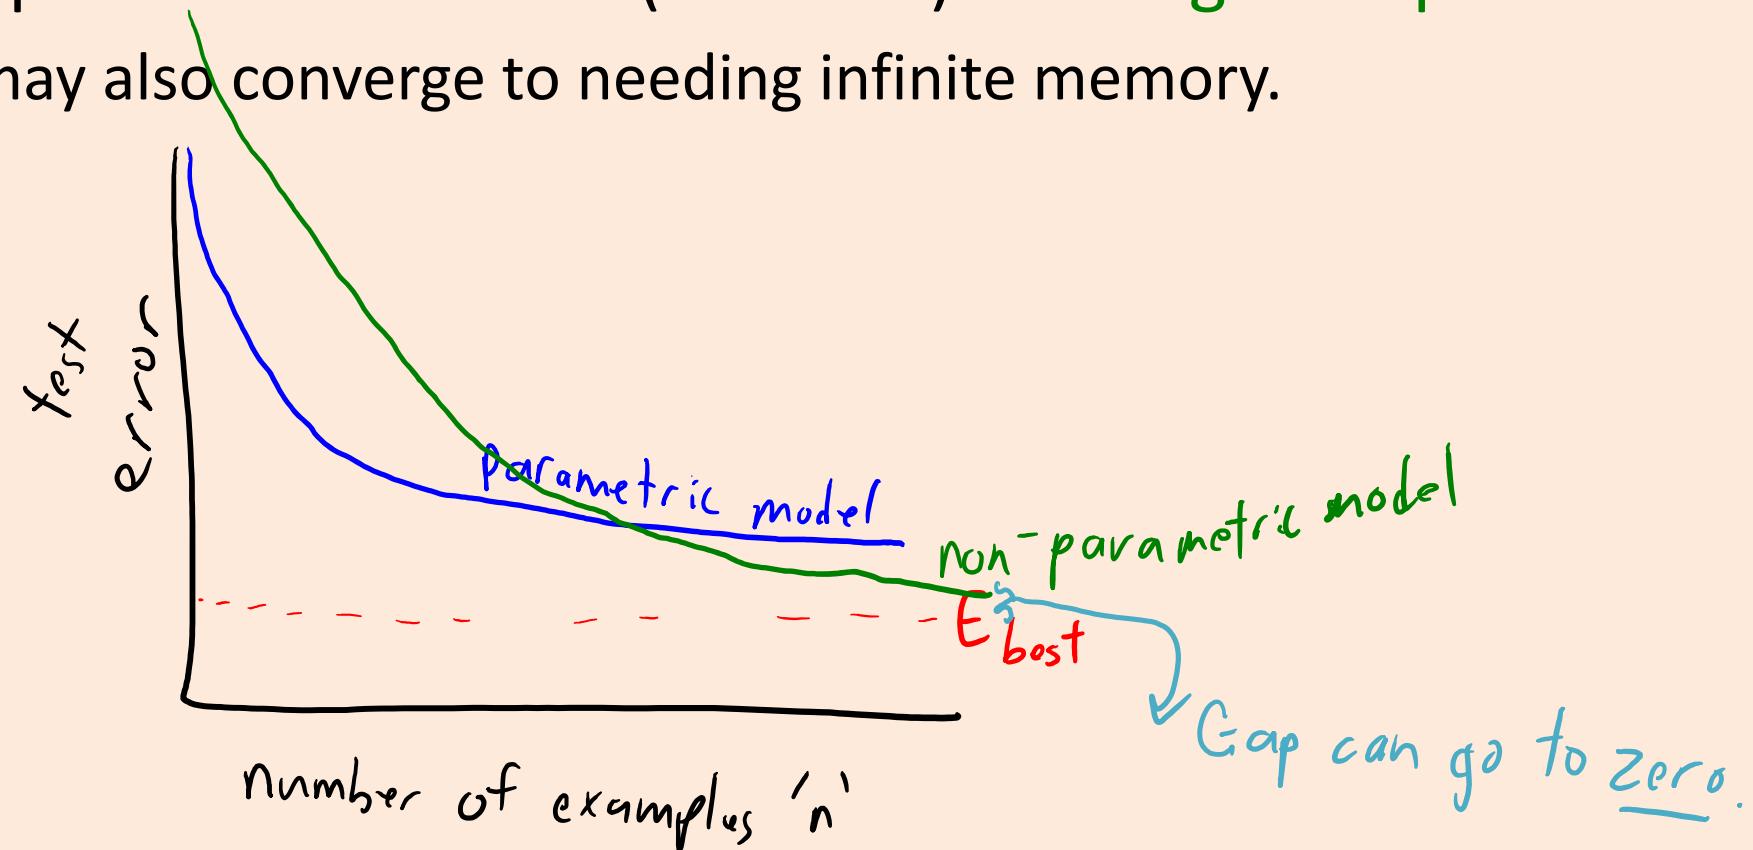
- With parametric models, there is an **accuracy limit**.
 - Even with infinite 'n', may not be able to achieve optimal error (E_{best}).



bonus!

Parametric vs. Non-Parametric Models

- With parametric models, there is an **accuracy limit**.
 - Even with infinite 'n', may not be able to achieve optimal error (E_{best}).
- Many non-parametric models (like kNN) **converge to optimal error**.
 - Though may also converge to needing infinite memory.



Summary

- **Decision theory** allows us to consider costs of predictions.
- **K-Nearest Neighbours:** use most common label of nearest examples.
 - Often works surprisingly well.
 - Suffers from high prediction and memory cost.
 - Canonical example of a “non-parametric” model.
 - Can suffer from the “curse of dimensionality”.
- **Non-parametric models** grow with number of training examples.
 - Can have appealing “consistency” properties (test error goes down to smallest possible error the model can make, as n goes to infinity).
- Next Time:
 - Fighting the fundamental trade-off.

Naïve Bayes Training Phase

- Training a naïve Bayes model:

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 1 \\ 1 & 0 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Naïve Bayes Training Phase

- Training a naïve Bayes model:

1. Set n_c to the number of times ($y_i = c$).

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 1 \\ 1 & 0 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$n_1 = 6$
 $n_0 = 4$

Naïve Bayes Training Phase

$$p(y_i=1) = \frac{6}{10} \leftarrow n_1 = 6$$

- Training a naïve Bayes model:

1. Set n_c to the number of times ($y_i = c$).

2. Estimate $p(y_i = c)$ as $\frac{n_c}{n}$.

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 1 \\ 1 & 0 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$p(y_i=0) = \frac{4}{10} \leftarrow n_0 = 4$$

Naïve Bayes Training Phase

- Training a naïve Bayes model:

1. Set n_c to the number of times ($y_i = c$).

2. Estimate $p(y_i = c)$ as $\frac{n_c}{n}$.

3. Set n_{cjk} as the number of times ($y_i = c, x_{ij} = k$)

$$X = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Red annotations:

- $p(y_i=1) = \frac{6}{10} \leftarrow n_1 = 6$ (highlighting the first column of the data matrix)
- $n_{121} = 4$ (highlighting the count of '1's in the first column of the data matrix)
- $p(y_i=0) = \frac{4}{10} \leftarrow n_0 = 4$ (highlighting the count of '0's in the first column of the data matrix)

Naïve Bayes Training Phase

$$p(y_i=1) = \frac{6}{10} \leftarrow n_1 = 6$$

- Training a naïve Bayes model:

1. Set n_c to the number of times ($y_i = c$).

2. Estimate $p(y_i = c)$ as $\frac{n_c}{n}$.

3. Set n_{cjk} as the number of times ($y_i = c, x_{ij} = k$)

4. Estimate $p(x_{ij} = k, y_i = c)$ as $\frac{n_{cjk}}{n}$.

$X =$	$\begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{bmatrix}, y =$	$\begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$
		$n_{121} = 4$

$$p(x_{12} = 1, y_i = 1) = \frac{4}{10}$$

$$p(y_i = 0) = \frac{4}{10} \leftarrow n_0 = 4$$

Naïve Bayes Training Phase

- Training a naïve Bayes model:

1. Set n_c to the number of times ($y_i = c$).

2. Estimate $p(y_i = c)$ as $\frac{n_c}{n}$.

3. Set n_{cjk} as the number of times ($y_i = c, x_{ij} = k$)

4. Estimate $p(x_{ij} = k, y_i = c)$ as $\frac{n_{cjk}}{n}$.

5. Use that $p(x_{ij} = k | y_i = c) = \frac{p(x_{ij} = k, y_i = c)}{p(y_i = c)}$

$$= \frac{n_{cjk}/n}{n_c/n} = \frac{n_{cjk}}{n_c}$$

$p(x_{i2} = 1 | y_i = 1) = \frac{4}{6} = \frac{2}{3}$

$p(y_i = 0) = \frac{4}{10} \leftarrow n_0 = 4$

$p(y_i = 1) = \frac{6}{10} \leftarrow n_1 = 6$

$$X = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{bmatrix}, y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$n_{121} = 4$

$p(x_{i2} = 1, y_i = 1) = \frac{4}{10}$

Naïve Bayes Prediction Phase

- Prediction in a naïve Bayes model:

Given a test example \tilde{x}_i we set prediction \tilde{y}_i^1 to the 'c' maximizing $p(\tilde{x}_i | \tilde{y}_i = c)$

Under the naïve Bayes assumption we can maximize:

$$p(\tilde{y}_i = c | \tilde{x}_i) \propto \prod_{j=1}^d [p(\tilde{x}_{ij} | \tilde{y}_i = c)] p(\tilde{y}_i = c)$$

Naïve Bayes Prediction Phase

- Prediction in a naïve Bayes model:

Consider $\tilde{x}_i = [1 \ 1]$ in this data set \longrightarrow

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 1 \\ 1 & 0 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Naïve Bayes Prediction Phase

- Prediction in a naïve Bayes model:

Consider $\tilde{x}_i = [1 \ 1]$ in this data set \rightarrow

$$p(\tilde{y}_i=0 | \tilde{x}_i) \propto p(\tilde{x}_{i1}=1 | \tilde{y}_i=0) p(\tilde{x}_{i2}=1 | \tilde{y}_i=0) p(\tilde{y}_i=0)$$
$$= (1) \quad (0.25) \quad (0.4) = 0.1$$

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 1 \\ 1 & 0 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Naïve Bayes Prediction Phase

- Prediction in a naïve Bayes model:

Consider $\tilde{x}_i = [1 \ 1]$ in this data set \rightarrow

$$p(\tilde{y}_i=0 | \tilde{x}_i) \propto p(\tilde{x}_{i1}=1 | \tilde{y}_i=0) p(\tilde{x}_{i2}=1 | \tilde{y}_i=0) p(\tilde{y}_i=0) \\ = (1) \quad (0.25) \quad (0.4) = 0.1$$

$$p(\tilde{y}_i=1 | \tilde{x}_i) \propto p(\tilde{x}_{i1}=1 | \tilde{y}_i=1) p(\tilde{x}_{i2}=1 | \tilde{y}_i=1) p(\tilde{y}_i=1) \\ = (0.5) \quad (0.666\dots) \quad (0.6) = 0.2$$

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 1 \\ 1 & 0 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Naïve Bayes Prediction Phase

- Prediction in a naïve Bayes model:

Consider $\tilde{x}_i = [1 \ 1]$ in this data set \rightarrow

$$p(\tilde{y}_i=0 | \tilde{x}_i) \propto p(\tilde{x}_{i1}=1 | \tilde{y}_i=0) p(\tilde{x}_{i2}=1 | \tilde{y}_i=0) p(\tilde{y}_i=0)$$
$$= (1) \quad (0.25) \quad (0.4) = 0.1$$

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 1 \\ 1 & 0 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$p(\tilde{y}_i=1 | \tilde{x}_i) \propto p(\tilde{x}_{i1}=1 | \tilde{y}_i=1) p(\tilde{x}_{i2}=1 | \tilde{y}_i=1) p(\tilde{y}_i=1)$$
$$= (0.5) \quad (0.666\dots) \quad (0.6) = 0.2$$

Since $p(\tilde{y}_i=1 | \tilde{x}_i)$ is bigger than $p(\tilde{y}_i=0 | \tilde{x}_i)$, naïve Bayes predicts $\hat{y}_i=1$.

(Don't sum to 1 because we're ignoring $p(\tilde{x}_i)$)

“Proportional to” for Probabilities

bonus!

- When we say “ $p(y) \propto \exp(-y^2)$ ” for a function ‘ p ’, we mean:

$$p(y) = \beta \exp(-y^2) \text{ for some constant } \beta.$$

- However, if ‘ p ’ is a probability then it must sum to 1.

- If $y \in \{1, 2, 3, 4\}$ then $p(1) + p(2) + p(3) + p(4) = 1$

- Using this fact, we can find β :

$$\beta \exp(-1^2) + \beta \exp(-2^2) + \beta \exp(-3^2) + \beta \exp(-4^2) = 1$$
$$\Leftrightarrow \beta [\exp(-1^2) + \exp(-2^2) + \exp(-3^2) + \exp(-4^2)] = 1$$

$$\Leftrightarrow \beta = \frac{\exp(-1^2) + \exp(-2^2) + \exp(-3^2) + \exp(-4^2)}{1}$$

bonus!

Probability of Paying Back a Loan and Ethics

- Article discussing predicting “whether someone will pay back a loan”:
 - <https://www.thecut.com/2017/05/what-the-words-you-use-in-a-loan-application-reveal.html>
- Words that **increase probability** of paying back the most:
 - *debt-free, lower interest rate, after-tax, minimum payment, graduate.*
- Words that **decrease probability** of paying back the most:
 - *God, promise, will pay, thank you, hospital.*
- Article also discusses an important issue: **are all these features ethical?**
 - Should you deny a loan because of religion or a family member in the hospital?
 - ICBC is limited in the features it is allowed to use for prediction.

bonus!

Avoiding Underflow

- During the prediction, the **probability** can underflow:

$$p(y_i=c | x_i) \propto \prod_{j=1}^d [p(x_{ij} | y_i=c)] p(y_i=c)$$

All these are < 1 so the product gets very small.

- Standard fix is to (equivalently) maximize the **logarithm** of the probability:

Remember that $\log(ab) = \log(a) + \log(b)$ so $\log(\prod a_i) = \sum_i \log(a_i)$

Since \log is monotonic the 'c' maximizing $p(y_i=c | x_i)$ also maximizes $\log p(y_i=c | x_i)$,

so maximize $\log \left(\prod_{j=1}^d [p(x_{ij} | y_i=c)] p(y_i=c) \right) = \sum_{j=1}^d \log(p(x_{ij} | y_i=c)) + \log(p(y_i=c))$

bonus!

Less-Naïve Bayes

- Given features $\{x_1, x_2, x_3, \dots, x_d\}$, naïve Bayes approximates $p(y|x)$ as:

$$\begin{aligned} p(y|x_1, x_2, \dots, x_d) &\propto p(y) p(x_1, x_2, \dots, x_d | y) && \downarrow \text{product rule applied repeatedly} \\ &= p(y) p(x_1 | y) p(x_2 | x_1, y) p(x_3 | x_2, x_1, y) \cdots p(x_d | x_1, x_2, \dots, x_{d-1}, y) \\ &\approx p(y) p(x_1 | y) p(x_2 | y) p(x_3 | y) \cdots p(x_d | y) && (\text{naive Bayes assumption}) \end{aligned}$$

- The assumption is very strong, and there are “less naïve” versions:

- Assume independence of all variables except up to ‘k’ largest ‘j’ where $j < i$.

- E.g., naïve Bayes has $k=0$ and with $k=2$ we would have:

$$\approx p(y) p(x_1 | y) p(x_2 | x_1, y) p(x_3 | x_2, x_1, y) p(x_4 | x_3, x_2, y) \cdots p(x_d | x_{d-2}, x_{d-1}, y)$$

- Fewer independence assumptions so more flexible, but hard to estimate for large ‘k’.
 - Another practical variation is “tree-augmented” naïve Bayes.

Computing $p(x_i)$ under naïve Bayes

bonus!

- Generative models don't need $p(x_i)$ to make decisions.
- However, it's easy to calculate under the naïve Bayes assumption:

$$p(x_i) = \sum_{c=1}^K p(x_i, y=c) \quad (\text{marginalization rule})$$

$$= \sum_{c=1}^K p(x_i | y=c) p(y=c) \quad (\text{product rule})$$

$$= \sum_{c=1}^K \left[\prod_{j=1}^d p(x_{ij} | y=c) \right] p(y=c) \quad (\text{naïve Bayes assumption})$$

These are the quantities
we compute during training.

bonus!

Gaussian Discriminant Analysis

- Classifiers based on Bayes rule are called **generative classifier**:
 - They often work well when you have **tons of features**.
 - But they **need to know $p(x_i | y_i)$** , probability of features given the class.
 - How to “generate” features, based on the class label.
- To fit generative models, usually make **BIG assumptions**:
 - **Naïve Bayes (NB)** for discrete x_i :
 - Assume that each variables in x_i is **independent** of the others in x_i given y_i .
 - **Gaussian discriminant analysis (GDA)** for continuous x_i .
 - Assume that $p(x_i | y_i)$ follows a multivariate normal distribution.
 - If all classes have same covariance, it’s called “linear discriminant analysis”.

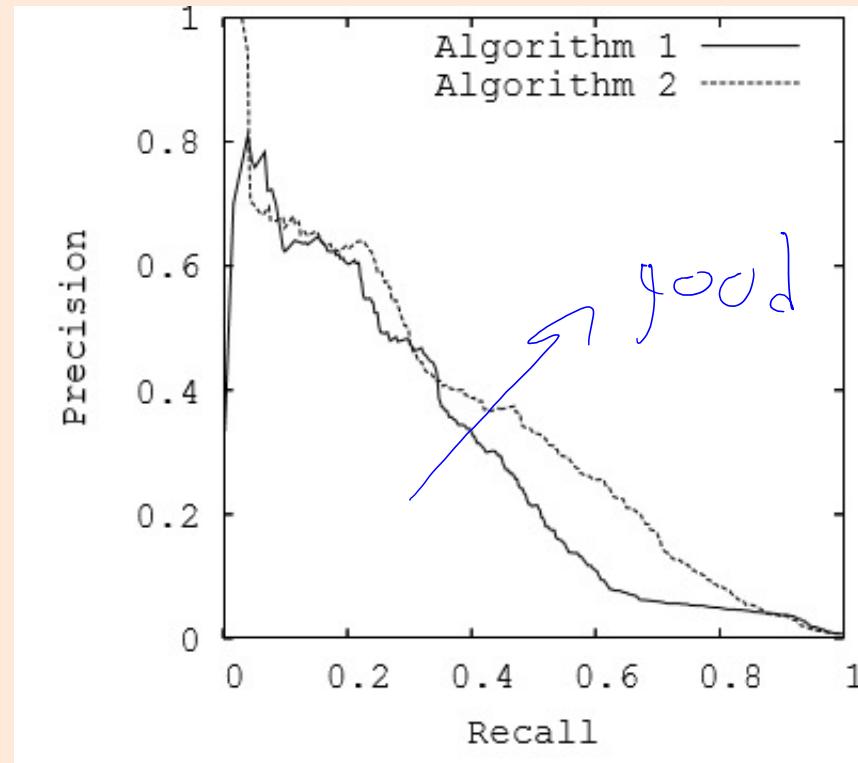
Other Performance Measures

- Classification error might be wrong measure:
 - Use weighted classification error if have different costs.
 - Might want to use things like Jaccard measure: $TP/(TP + FP + FN)$.
- Often, we report **precision** and **recall** (**want both to be high**):
 - Precision: “if I classify as spam, what is the probability it actually is spam?”
 - Precision = $TP/(TP + FP)$.
 - High precision means the filtered messages are likely to really be spam.
 - Recall: “if a message is spam, what is probability it is classified as spam?”
 - Recall = $TP/(TP + FN)$
 - High recall means that most spam messages are filtered.

bonus!

Precision-Recall Curve

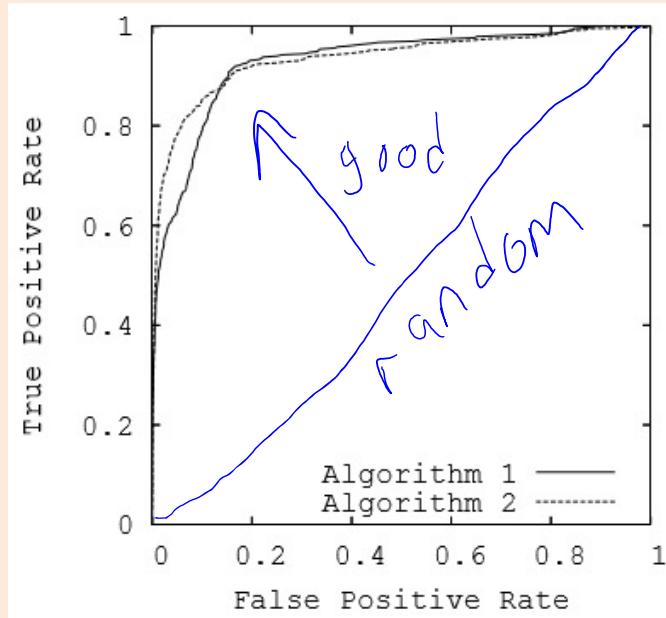
- Consider the rule $p(y_i = \text{'spam'} | x_i) > t$, for threshold 't'.
- Precision-recall (PR) curve plots precision vs. recall as 't' varies.



bonus!

ROC Curve

- Receiver operating characteristic (ROC) curve:
 - Plot true positive rate (recall) vs. false positive rate ($FP/FP+TN$).
(negative examples classified as positive)



- Diagonal is random, perfect classifier would be in upper left.
- Sometimes papers report area under curve (AUC).
 - Reflects performance for different possible thresholds on the probability.

bonus!

More on Unbalanced Classes

- With unbalanced classes, there are many alternatives to accuracy as a measure of performance:
 - Two common ones are the Jaccard coefficient and the F-score.
- Some machine learning models don't work well with unbalanced data. Some common heuristics to improve performance are:
 - Under-sample the majority class (only take 5% of the spam messages).
 - <https://www.jair.org/media/953/live-953-2037-jair.pdf>
 - Re-weight the examples in the accuracy measure (multiply training error of getting non-spam messages wrong by 10).
 - Some notes on this issue are [here](#).

bonus!

More on Weirdness of High Dimensions

- In high dimensions:
 - Distances become less meaningful:
 - All vectors may have similar distances.
 - Emergence of “hubs” (even with random data):
 - Some datapoints are neighbours to many more points than average.
 - Visualizing high dimensions and sphere-packing

bonus!

Vectorized Distance Calculation

- To classify ‘t’ test examples based on kNN, cost is $O(ndt)$.
 - Need to compare ‘n’ training examples to ‘t’ test examples, and computing a distance between two examples costs $O(d)$.
- You can do this using matrix multiplication:
 - Let D be a matrix such that D_{ij} contains:

$$\|x_i - x_j\|^2 = \|x_i\|^2 - 2x_i^T x_j + \|x_j\|^2$$

where ‘i’ is a training example and ‘j’ is a test example.

- In numpy: (like [sklearn.metrics.pairwise.euclidean_distances](#))

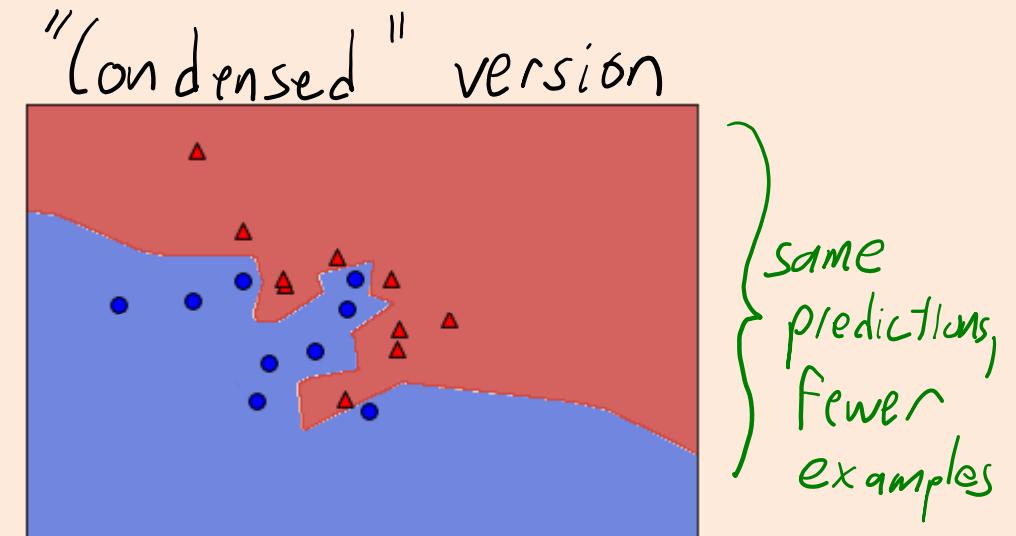
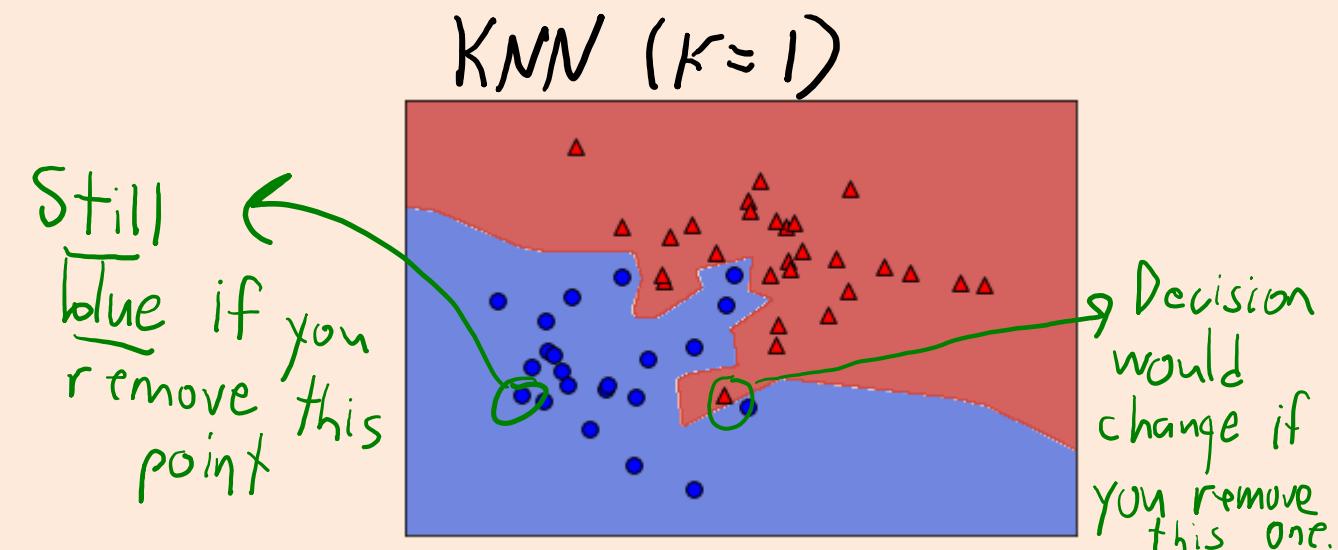


- Can be better than optimized C loops ([scipy.spatial.distance.cdist](#))

bonus!

Condensed Nearest Neighbours

- Disadvantage of kNN is **slow prediction time** (depending on 'n').
- **Condensed nearest neighbours:**
 - Identify a set of 'm' "prototype" training examples.
 - Make predictions by using these "prototypes" as the training data.
- Reduces runtime from $O(nd)$ down to $O(md)$.



bonus!

Condensed Nearest Neighbours

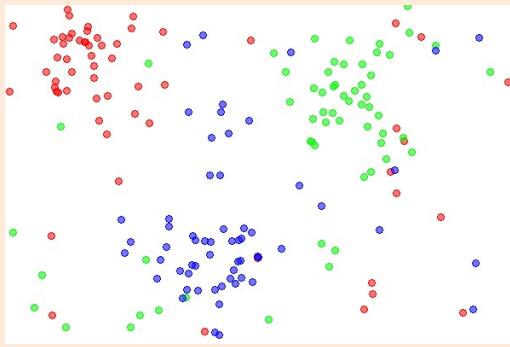
- Classic condensed nearest neighbours:
 - Start with no examples among prototypes.
 - Loop through the non-prototype examples ‘i’ in some order:
 - Classify x_i based on the current prototypes.
 - If prediction is not the true y_i , add it to the prototypes.
 - Repeat the above loop until all examples are classified correctly.
- Some variants first remove points from the original data, if a full-data KNN classifier classifies them incorrectly (“outliers”).

bonus!

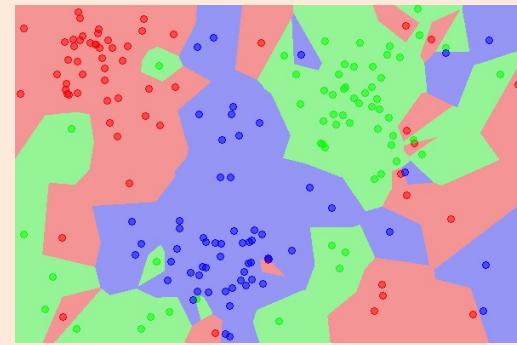
Condensed Nearest Neighbours

- Classic condensed nearest neighbours:

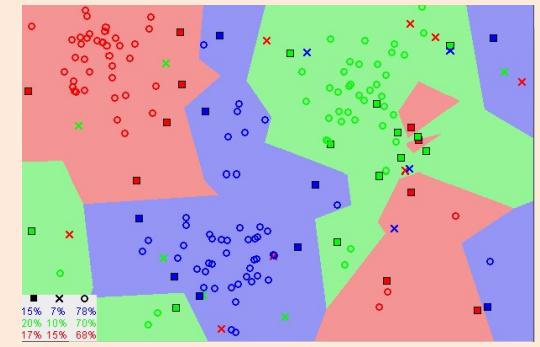
Data:



1NN



CNN:
□: prototype
○: removed
×: outlier
(using 3NN)



- Recent work shows that finding optimal compression is NP-hard.
 - An approximation algorithm was published in 2018:
 - [“Near optimal sample compression for nearest neighbors”](#)

bonus!

Approximate Nearest Neighbours

- Store data in a special data structure, e.g. k-d tree
 - Partition points into regions, only check nearby regions
 - Only helps for exact checks if n is at least about 2^d
 - But making several trees on different projections can give good approximations (that might miss true NNs)



from [vlfeat docs](#)

- Locality-sensitive hashing
 - Like traditional hashing but we *try to get collisions* for nearby points
 - Simple method (SimHash): choose random hyperplanes, track which side of each the result is on

bonus!

Refined Fundamental Trade-Off

- Let E_{best} be the **irreducible error** (lowest possible error for *any* model).
 - For example, irreducible error for predicting coin flips is 0.5.
- Some learning theory results use E_{best} to further decompose E_{test} :

$$E_{\text{test}} = \underbrace{(E_{\text{test}} - E_{\text{train}})}_{E_{\text{approx}}} + \underbrace{(E_{\text{train}} - E_{\text{best}})}_{E_{\text{model}}} + \underbrace{E_{\text{best}}}_{\text{"noise"}}$$

- E_{approx} measures *how sensitive we are to training data*.
- E_{model} measures *if our model is complicated enough to fit data*.
- E_{best} measures how low can **any** model make test error.
 - E_{best} does not depend on what model you choose.

Consistency and Universal Consistency

- A model is **consistent** for a **particular learning problem** if:
 - E_{test} converges to E_{best} as ‘n’ goes to infinity, for that particular problem.
- A model is **universally consistent** for a **class of learning problems** if:
 - E_{test} converges to E_{best} as ‘n’ goes to infinity, for all problems in the class.
- **Class of learning problems** will usually be “all problems satisfying”:
 - A **continuity assumption** on the labels y^i as a function of x^i .
 - E.g., if x^i is close to x^j then they are likely to receive the same label.
 - A **boundedness assumption** of the set of x^i .

bonus!

Consistency of KNN (Discrete/Deterministic Case)

- Let's show universal consistency of KNN in a simplified setting.
 - The x^i and y^i are binary, and y^i being a deterministic function of x^i .
 - Deterministic y^i implies that E_{best} is 0.
- Consider KNN with $k=1$:
 - After we observe an x_i , KNN makes right test prediction for that vector.
 - As 'n' goes to ∞ , each feature vectors with non-zero probability is observed.
 - We have $E_{\text{test}} = 0$ once we've seen all feature vectors with non-zero probability.
- Notes:
 - “No free lunch” isn’t relevant as ‘n’ goes to ∞ : we eventually see everything.
 - But there are 2^d possible feature vectors, so might need a huge number of training examples.
 - It’s more complicated if labels aren’t deterministic and features are continuous.

bonus!

Consistency of Non-Parametric Models

- **Universal consistency** can be shown for many models we'll cover:
 - Linear models with polynomial basis.
 - Linear models with Gaussian RBFs.
 - Neural networks with one hidden layer and standard activations.
 - Sigmoid, tanh, ReLU, etc.
- But it's always the **non-parametric versions** that are consistent:
 - Where **size of model** is a function of 'n'.
 - Examples:
 - KNN needs to store all 'n' training examples.
 - Degree of polynomial must grow with 'n' (not true for fixed polynomial).
 - Number of hidden units must grow with 'n' (not true for fixed neural network).